Siddhartha Reddy Kothakapu

http://www.grok.in/
http://www.linkedin.com/in/sidsr

sids@sids.in
+91 998 028 7437

EXPERIENCE

Technical Guru, Ziva Software, Bangalore   *Jul 2006 - Present*

Ziva Software is a startup based out of Bangalore, providing products and services in the *Mobile Search* domain. Established in 2005, Ziva has provided search solutions to several players in the mobile industry including operators, manufacturers and content owners. In 2007, Ziva launched *Zook*, its consumer facing mobile search engine.

Was the first non-founding employee. Have been a core member of the team and was involved in most of the projects. Helped define both development and strategic roadmap for the company.

EDUCATION

M.Tech in Information Technology (3.71/4.00 CGPA)   *Jun 2006*

International Institute of Information Technology, Bangalore

SKILLS

Languages: Perl, Java, PHP, Ruby, Javascript, C/C++,

Frameworks/Libraries/Softwares: Hadoop, Nutch, Lucene, Solr, Sphinx, WEKA, Ruby on Rails, jQuery, MySQL

Others: HTML, CSS, JSP, Bash Shell Scripting

INTERESTS

Information Engineering—Retrieval, Extraction & Management, Machine Learning, Scalable Architectures, Distributed & Cloud Computing, Usability & User Experience

SELECTED
PROJECTS

Buddhi Online, Ziva Software   TEAM SIZE: 1-2   *Nov 2008 - Present*

Real-time extraction of actionable *information snippets* from web pages to directly provide exact answers to search queries.

Conceptualised, designed and developed the complete system — first as a prototype and then as a production system that is currently powering Zook.

Buddhi, Ziva Software   TEAM SIZE: 3-4   *Apr 2007 - Present*

A dream to crawl the Web and automagically extract *information-content entities* off the crawled pages. That is, Buddhi extracts (semi-)structured data out of *any* web page without any apriori information about it.

Have been part of the project right from its inception and have been managing it for the past few months.

⋄ Developed the initial prototype for a feasibility study.
⋄ Designed the *Map-Reduce* based architecture for the system.
⋄ Implemented a *focused crawler* based on the *Nutch* open source crawler. The crawler has been deployed and regularly crawls large portions of the Web.
⋄ Set up the *Hadoop* based infrastructure for the crawling and extraction.

Zook, Ziva Software   TEAM SIZE: 2-3   *Jan 2007 - Present*

Ziva Software's flagship mobile search engine. Zook's 1-Click Search provides actionable information nuggets instead of a set of links to websites that *might* have that information.

Have been involved in all aspects of Zook's development and deployment.

⋄ Set up the search system infrastructure.
⋄ Designed the *UI* and *UX* of the WAP and SMS interfaces.
⋄ Responsible for deploying releases.
⋄ Created a HTTP API for accessing the search system.

Zivabot, Ziva Software   Team Size: 1   *Aug 2006 - Oct 2006*

A scalable, highly-configurable and generic *scraping system* for crawling specific websites containing desired information and then extracting that information into a pre-defined schema.

Was responsible for the design and implementation of the system. Configured the system and set it up to work with an initial set of websites; this provided the base content for the launch of Zook.

Better Web Buys   Team Size: 2   *Feb 2006 - May 2006*

A *faceted price comparison shopping search engine* with over 20 million items across more than 100 merchants.

Was completely responsible for the implementation right from fetching the product feeds, parsing and storing them in a database, indexing, searching, user experience and setting up the system on a cluster of machines. Used *Sphinx* as the backend full-text indexing system.

yadb – Yet Another DataBase, IIIT-Bangalore   Team Size: 3   *Aug 2005 - Dec 2005*

A reasonably *complete RDBMS* with a buffered memory manager, B+ tree indexing and a complete database management module including support for basic SQL, table management, paged storage etc. Part of the *Advanced Databases* course.

Top News, IIIT-Bangalore   Team Size: 1   *Sep 2005 - Oct 2005*

A system to bring out the top news of the moment by *clustering the news items* obtained from the feeds provided by several news sources. Any news item being discussed by multiple sources within a given (variable) timeframe is marked as top news. Final project for the *Web Information Retreival* course.

SiDuSE, IIIT-Bangalore   Team Size: 2   *Aug 2004 - Sep 2004*

A *Simple Desktop Search Engine* for Linux/UNIX systems. Creates a Trie-based full-text index of the files on the system. Apart from allowing the basic *AND*, *OR* & *NOT* operators, the system supports complex querying through nested queries, phrase matching and prefix queries. Part of the *Mathematical Foundations of Computing* course.

SOS – S Operating System   Team Size: 2   *Sep 2002 - Oct 2003*

A simple monolithic kernel with kernel level threads, kernel level IPC and advanced memory management.

Publication   **Measures of "Ignorance" on the Web**   *Dec 2006*

International Conference on Management of Data, Delhi, India

Honors

⋄ *Valedictorian* for the class for 2006 at IIIT-Bangalore.

⋄ Awarded the *Honeywell Scholarship* for academic excellence at IIIT-Bangalore.

⋄ Chosen as the *employee of the year* for 2007-08 at Ziva Software.

⋄ Adjudged the *best "Really Needs an Interface"* hack at the Yahoo! Open Hackday India 2007.

⋄ Elected the president of the IEEE student chapter at the J.B. Institute of Engineering & Technology. Was responsible for organising two of the institutes technical festivals in this position.